

Average Number of Accepted Tokens per Step

LLM: LLAMA-3.1-70B-Instruct

Batch Sizes: 4, 8

